

# Prosperity Prognosticator

## Machine Learning for Startup Success Prediction

### Project Description:

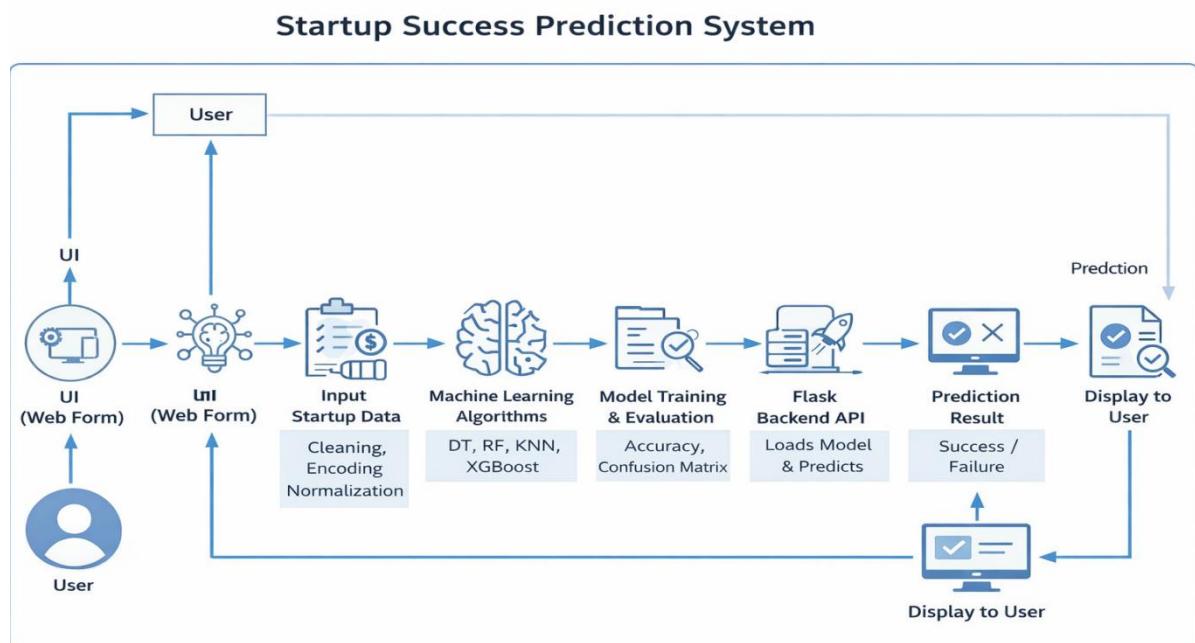
One of the most important factors influencing a country's economic growth and innovation ecosystem is the success and sustainability of startups. Startups play a crucial role in job creation, technological advancement, and economic development. However, predicting whether a startup will succeed or fail is a challenging task due to multiple influencing factors such as funding, market trends, team experience, and industry growth.

Startup success prediction is a complex problem, and many traditional evaluation methods fail to capture the hidden patterns in large datasets. By forecasting startup success or failure, investors, entrepreneurs, and policymakers can make better decisions, reduce financial risks, and allocate resources effectively.

In this project, we will use machine learning classification and regression algorithms such as Logistic Regression, Decision Tree, Random Forest, KNN, and XGBoost to predict startup success. The dataset will be preprocessed using data cleaning, feature engineering, and exploratory data analysis techniques. Finally, the model will be integrated into a Flask-based web application and deployed on a cloud platform (such as IBM Cloud) to provide real-time startup success predictions.

This system will help investors identify promising startups, assist entrepreneurs in improving their business strategies, and support policymakers in designing programs that promote innovation and economic development.

### Technical Architecture:



## **Pre-requisites**

To successfully complete the **Startup Success Prediction System** project, the following software tools, concepts, and Python packages are required.

### **1. Software Requirements**

#### **Anaconda Navigator and PyCharm**

Anaconda is used to manage Python environments and install required packages easily. PyCharm is used as the Integrated Development Environment (IDE) for coding and testing the project.

##### **Download Anaconda Navigator:**

You can refer to the following tutorial link for installation:

<https://youtu.be/1ra4zH2G4o0>

- **Anaconda navigator and pycharm:**
  - Refer the link below to download anaconda navigator
  - Link : [https://youtu.be/mg6cMkz9Q0c?si=O\\_e6\\_4qgGF8S-lut](https://youtu.be/mg6cMkz9Q0c?si=O_e6_4qgGF8S-lut)

#### **Python Concepts Required**

- Python Programming Basics
- Exploratory Data Analysis (EDA)
- Data Preprocessing Techniques
- Machine Learning Concepts
- Classification Algorithms
- Model Evaluation Techniques

#### **Required Python Packages**

Open **Anaconda Prompt as Administrator** and install the following packages using the commands below:

pip install numpy

pip install pandas

pip install scikit-learn

pip install matplotlib

pip install scipy

pip install pickle-mixin

pip install seaborn

pip install Flask

Flask Web Application Development

## Prior Knowledge:

You must have prior knowledge of following topics to complete this project.

- **ML Concepts**
  - Supervised learning: <https://www.javatpoint.com/supervised-machine-learning>
  - Unsupervised learning: <https://www.javatpoint.com/unsupervised-machine-learning>
  - Regression and classification
  - Decision tree: <https://www.javatpoint.com/machine-learning-decision-tree-classification-algorithm>
  - Random forest: <https://www.javatpoint.com/machine-learning-random-forest-algorithm>
  - KNN: <https://www.javatpoint.com/k-nearest-neighbor-algorithm-for-machine-learning>
  - Xgboost: <https://www.analyticsvidhya.com/blog/2018/09/an-end-to-end-guide-to-understand-the-math-behind-xgboost/>
  - Evaluation metrics: <https://www.analyticsvidhya.com/blog/2019/08/11-important-model-evaluation-error-metrics/>
- **Flask Basics** : [https://www.youtube.com/watch?v=lj4l\\_CvBnt0](https://www.youtube.com/watch?v=lj4l_CvBnt0)

## Project Objectives:

By the end of this project you will:

- Know fundamental concepts and techniques used for machine learning.
- Gain a broad understanding about data.
- Have knowledge on pre-processing the data/transformation techniques on outlier and some visualization concepts.

## Project Flow:

- User interacts with the UI to enter the input.
- Entered input is analyzed by the model which is integrated.
- Once model analyses the input the prediction is showcased on the UI

To accomplish this, we have to complete all the activities listed below,

- Data collection
  - Collect the dataset or create the dataset
- Visualizing and analyzing data
  - Univariate analysis
  - Bivariate analysis
  - Multivariate analysis
  - Descriptive analysis
- Data pre-processing
  - Checking for null values
  - Handling outlier
  - Handling categorical data
  - Splitting data into train and test
- Model building
  - Import the model building libraries
  - Initializing the model
  - Training and testing the model
  - Evaluating performance of model
  - Save the model
- Application Building
  - Create an HTML file
  - Build python code

## Project Structure:

Create the Project folder which contains files as shown below

Name	Date modified	Type	Size
▼ Today			
app.py	2/18/2026 2:45 PM	Python Source File	2 KB
random_forest_model.pkl	2/18/2026 2:45 PM	PKL File	1,574 KB
Startup_Success_Prediction (1).ipynb	2/18/2026 2:45 PM	IPYNB File	2,482 KB
startup1.csv	2/18/2026 2:45 PM	XLS Worksheet	206 KB
template	2/18/2026 2:45 PM	File folder	

- We are building a flask application which needs HTML pages stored in the templates folder and a python script app.py for scripting.
- Random\_forest\_model.pkl is our saved model. Further we will use this model for flask integration.
- Training folder contains model training files and training\_ibm folder contains IBM deployment files.

## Milestone 1: Data Collection

ML depends heavily on data, It is most crucial aspect that makes algorithm training possible. So this section allows you to download the required dataset.

### Activity 1: Download the dataset

There are many popular open sources for collecting the data. Eg: kaggle.com

In this project we have used startup1.csv data. This data is downloaded from kaggle.com. Please refer to the link given below to download the dataset.

Link: [startup1.csv](#)

## Milestone 2: Data Information and Initial Exploration

As the dataset is downloaded. Let us read and understand the data properly with the help of some visualization techniques and some analysing techniques.

**Note: There is n number of techniques for understanding the data. But here we have used some of it. In an additional way, you can use multiple techniques.**

### Activity 1: Importing the libraries

Import the necessary libraries as shown in the image. (optional) Here we have used visualization style as fivethirtyeight.

```
#!pip install plotly
#!pip install pandas
#!pip install numpy
#pip install matplotlib
#pip install seaborn
✓import pandas as pd
import numpy as np
import matplotlib as mpl
from matplotlib import pyplot as plt
import seaborn as sns
from datetime import date
from scipy import stats
from scipy.stats import norm, skew #for some statistics
import warnings
warnings.filterwarnings("ignore")
```

## Activity 2: Read the Dataset:

Our dataset format might be in .csv, excel files, .txt, .json, etc. We can read the dataset with the help of pandas.

In pandas we have a function called `read_csv()` to read the dataset. As a parameter we have to give the directory of csv file.

Load Dataset

```
data= pd.read_csv('startup1.csv')
data
```

Python

	Unnamed: 0	state_code	latitude	longitude	zip_code	id	city	Unnamed: 6	name	labels	...	object_id	has_VC	has_angel	has_roundA	has_roundB	has_roundC	has_roundD	avg_participant
0	1005	CA	42.358880	-71.056820	92101	c6669	San Diego	NaN	Bandsintown	1	...	c6669	0	1	0	0	0	0	1.000
1	204	CA	37.238916	-121.973718	95032	c16283	Los Gatos	NaN	TriCipher	1	...	c16283	1	0	0	1	1	1	4.750
2	1001	CA	32.901049	-117.192656	92121	c65620	San Diego	San Diego CA 92121	Pixi	1	...	c65620	0	0	1	0	0	0	4.000
3	738	CA	37.320309	-122.050040	95014	c42668	Cupertino	Cupertino CA 95014	Solidcore Systems	1	...	c42668	0	0	0	1	1	1	3.333
4	1002	CA	37.779281	-122.419236	94105	c65806	San Francisco	San Francisco CA 94105	Inhale Digital	0	...	c65806	1	1	0	0	0	0	1.000
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
918	352	CA	37.740594	-122.376471	94107	c21343	San Francisco	NaN	CoTweet	1	...	c21343	0	0	1	0	0	0	6.000
919	721	MA	42.504817	-71.195611	1803	c41747	Burlington	Burlington MA 1803	Reef Point Systems	0	...	c41747	1	0	0	1	0	0	2.666
920	557	CA	37.408261	-122.015920	94089	c31549	Sunnyvale	NaN	Paracor Medical	0	...	c31549	0	0	0	0	0	1	8.000
921	589	CA	37.556732	-122.288378	94404	c33198	San Francisco	Santa Clara CA 95054	Causata	1	...	c33198	0	0	1	1	0	0	1.000
922	462	CA	37.386778	-121.966277	95054	c26702	Santa Clara	Santa Clara CA 95054	Asempra Technologies	1	...	c26702	0	0	0	1	0	0	3.000

923 rows × 49 columns

Before analysing, we need to understand what kind of data we have — the column names, data types, and null counts. This step guides all further preprocessing decisions.

## Activity 3: Data Info and Columns

`data.info()` is one of the most useful pandas functions. It shows each column name, its non-null count, and its data type (int64, float64, object). This immediately reveals which columns have missing values and which are numeric vs. categorical.

# Data Information

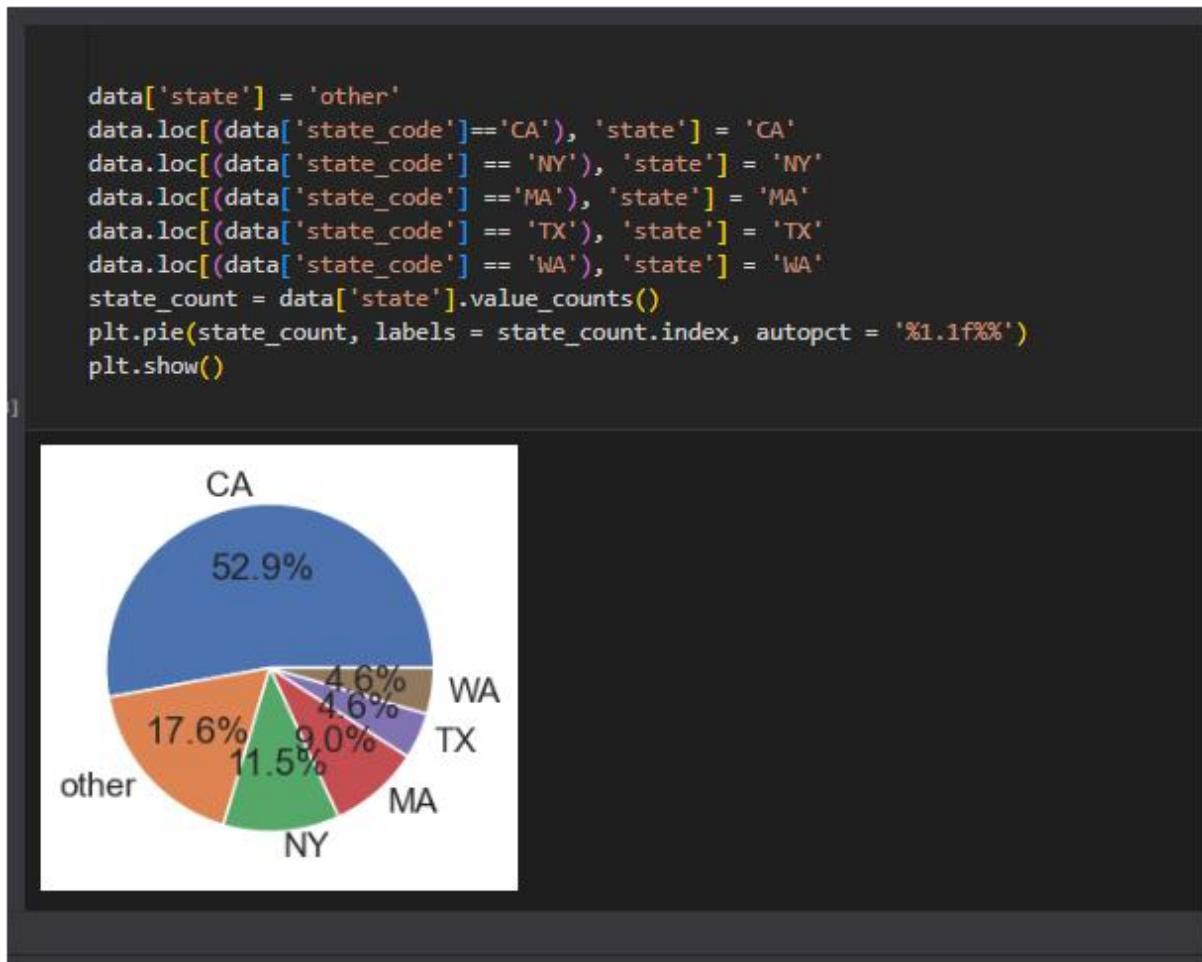
```
data.info()  
[136]  
... <class 'pandas.core.frame.DataFrame'>  
RangeIndex: 923 entries, 0 to 922  
Data columns (total 49 columns):  
 #   Column            Non-Null Count  Dtype     
---  --  
 0   Unnamed: 0          923 non-null    int64    
 1   state_code         923 non-null    object    
 2   latitude           923 non-null    float64   
 3   longitude          923 non-null    float64   
 4   zip_code           923 non-null    object    
 5   id                 923 non-null    object    
 6   city               923 non-null    object    
 7   Unnamed: 6          430 non-null    object    
 8   name               923 non-null    object    
 9   labels              923 non-null    int64    
 10  founded_at         923 non-null    object    
 11  closed_at          335 non-null    object    
 12  first_funding_at   923 non-null    object    
 13  last_funding_at    923 non-null    object    
 14  age_first_funding_year  923 non-null    float64   
 15  age_last_funding_year  923 non-null    float64   
 16  age_first_milestone_year  771 non-null    float64   
 17  age_last_milestone_year  771 non-null    float64   
 18  relationships       923 non-null    int64    
 19  funding_rounds      923 non-null    int64    
 ...  
 47  is_top500           923 non-null    int64    
 48  status              923 non-null    object    
dtypes: float64(7), int64(28), object(14)  
memory usage: 353.5+ KB  
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```

```
data.columns  
[137]  
... Index(['Unnamed: 0', 'state_code', 'latitude', 'longitude', 'zip_code', 'id',  
       'city', 'Unnamed: 6', 'name', 'labels', 'founded_at', 'closed_at',  
       'first_funding_at', 'last_funding_at', 'age_first_funding_year',  
       'age_last_funding_year', 'age_first_milestone_year',  
       'age_last_milestone_year', 'relationships', 'funding_rounds',  
       'funding_total_usd', 'milestones', 'state_code.1', 'is_CA', 'is_NY',  
       'is_MA', 'is_TX', 'is_otherstate', 'category_code', 'is_software',  
       'is_web', 'is_mobile', 'is_enterprise', 'is_advertising',  
       'is_gamesvideo', 'is_ecommerce', 'is_biotech', 'is_consulting',  
       'is_othercategory', 'object_id', 'has_VC', 'has_angel', 'has_roundA',  
       'has_roundB', 'has_roundC', 'has_roundD', 'avg_participants',  
       'is_top500', 'status'],  
       dtype='object')
```

## Activity 4: State Code Grouping and Pie Chart

The dataset contains many US state codes. To simplify analysis, we group the top 5 most frequent states (CA, NY, MA, TX, WA) and group all remaining states into a single 'other' category. We then visualize the distribution as a pie chart.

This helps us understand which geographies dominate the startup ecosystem and whether state location could be a useful feature for prediction.



## Activity 5: Separate Numeric and Categorical Data

Pandas select\_dtypes() lets us filter columns by their data type. This is important because numeric and categorical columns need very different handling — numerics can be scaled and used directly while categoricals need encoding or separate treatment.

### Numeric Data:

**Code:**

```

numeric = ['int8', 'int16', 'int32', 'int64', 'float16', 'float32',
'float64']

data_num = data.select_dtypes(include=numeric)

data_num.head()

```

Numeric data

The screenshot shows a Jupyter Notebook cell with the following code:

```

numeric=['int8', 'int16', 'int32', 'int64', 'float16', 'float32',
'data_num=data.select_dtypes(include=numeric)
data_num.head()']

```

The output displays the first five rows of the numeric data frame:

	Unnamed: 0	latitude	longitude	labels	age_first_funding_year	age_last_funding_year	age_first_milestone_year	age_last_milestone_year	relationships	funding_rounds	...	is_consulting	is_othercategory	has_VC
0	1005	42.358880	-71.056820	1	2.2493	3.0027	4.6685	6.7041	3	3	...	0	1	0
1	204	37.238916	-121.973718	1	5.1260	9.9973	7.0055	7.0055	9	4	...	0	0	1
2	1001	32.901049	-117.192656	1	1.0329	1.0329	1.4575	2.2055	5	1	...	0	0	0
3	738	37.320309	-122.050040	1	3.1315	5.3151	6.0027	6.0027	5	3	...	0	0	0
4	1002	37.779281	-122.419236	0	0.0000	1.6685	0.0384	0.0384	2	2	...	0	0	1

5 rows × 35 columns

## Categorical Data:

### Code:

```

data_cat = data.select_dtypes(include='object')

data_cat.head()

```

Categorical Data

The screenshot shows a Jupyter Notebook cell with the following code:

```

data_cat=data.select_dtypes(include='object')
data_cat.head()

```

The output displays the first five rows of the categorical data frame:

	state_code	zip_code	id	city	Unnamed: 6	name	founded_at	closed_at	first_funding_at	last_funding_at	state_code.1	category_code	object_id	status	state	
0	CA	92101	c6669	San Diego		NaN	Bandsintown	1/1/2007	NaN	4/1/2009	1/1/2010	CA	music	c6669	acquired	CA
1	CA	95032	c16283	Los Gatos		NaN	Tricipher	1/1/2000	NaN	2/14/2005	12/28/2009	CA	enterprise	c16283	acquired	CA
2	CA	92121	c65620	San Diego	San Diego CA 92121		Pixii	3/18/2009	NaN	3/30/2010	3/30/2010	CA	web	c65620	acquired	CA
3	CA	95014	c42668	Cupertino	Cupertino CA 95014	Solidcore Systems	1/1/2002	NaN	2/17/2005	4/25/2007	CA	software	c42668	acquired	CA	
4	CA	94105	c65806	San Francisco	San Francisco CA 94105	Inhale Digital	8/1/2010	10/1/2012	8/1/2010	4/1/2012	CA	games_video	c65806	closed	CA	

## Milestone 3: Visualizing and Analysing the Data

As the dataset is loaded and the target variable is prepared, we now understand the data deeply through statistical summaries and visual exploration. There are N number of techniques for understanding data — here we use several of the most effective ones.

### Activity 1: Statistical Analysis

The describe() function generates a statistical summary of the DataFrame columns. It shows count, mean, standard deviation, minimum, maximum, and percentile

values for numeric columns. This instantly highlights skew, spread, and outlier presence without a single graph.

## Numeric Statistics:

### Code:

```
de_Num = data.describe(include=['float64', 'int64', 'float', 'int'])
```

	count	mean	std	min	25%	50%	75%	max
Unnamed: 0	923.000000	572.297941	333.585431	1.000000	283.500000	577.000000	866.500000	1153.000000
latitude	923.000000	38.517442	3.741497	25.752358	37.388869	37.779281	40.730646	59.335232
longitude	923.000000	-103.539212	22.394167	-122.756956	-122.198732	-118.374037	-77.214731	18.057121
age_first_funding_year	923.000000	2.235630	2.510449	-9.046600	0.576700	1.446600	3.575350	21.895900
age_last_funding_year	923.000000	3.931456	2.967910	-9.046600	1.669850	3.528800	5.560250	21.895900
age_first_milestone_year	771.000000	3.055353	2.977057	-14.169900	1.000000	2.520500	4.686300	24.684900
age_last_milestone_year	771.000000	4.754423	3.212107	-7.005500	2.411000	4.476700	6.753400	24.684900
relationships	923.000000	7.710726	7.265776	0.000000	3.000000	5.000000	10.000000	63.000000
funding_rounds	923.000000	2.310943	1.390922	1.000000	1.000000	2.000000	3.000000	10.000000
funding_total_usd	923.000000	25419749.092091	189634564.488794	11000.000000	2725000.000000	10000000.000000	24725000.000000	5700000000.000000
milestones	923.000000	1.841820	1.322632	0.000000	1.000000	2.000000	3.000000	8.000000
is_CA	923.000000	0.527627	0.499507	0.000000	0.000000	1.000000	1.000000	1.000000
is_NY	923.000000	0.114843	0.319005	0.000000	0.000000	0.000000	0.000000	1.000000
is_MA	923.000000	0.089924	0.286228	0.000000	0.000000	0.000000	0.000000	1.000000
is_TX	923.000000	0.045504	0.208519	0.000000	0.000000	0.000000	0.000000	1.000000
is_otherstate	923.000000	0.221018	0.415158	0.000000	0.000000	0.000000	0.000000	1.000000
is_software	923.000000	0.165764	0.372070	0.000000	0.000000	0.000000	0.000000	1.000000
is_web	923.000000	0.156013	0.363064	0.000000	0.000000	0.000000	0.000000	1.000000
is_mobile	923.000000	0.085590	0.279910	0.000000	0.000000	0.000000	0.000000	1.000000
is_enterprise	923.000000	0.079090	0.270025	0.000000	0.000000	0.000000	0.000000	1.000000
is_advertising	923.000000	0.067172	0.250456	0.000000	0.000000	0.000000	0.000000	1.000000
is_gamesvideo	923.000000	0.056338	0.230698	0.000000	0.000000	0.000000	0.000000	1.000000
is_ecommerce	923.000000	0.027086	0.162421	0.000000	0.000000	0.000000	0.000000	1.000000
is_biotech	923.000000	0.036836	0.188462	0.000000	0.000000	0.000000	0.000000	1.000000
is_consulting	923.000000	0.003250	0.056949	0.000000	0.000000	0.000000	0.000000	1.000000
is_othercategory	923.000000	0.322860	0.467823	0.000000	0.000000	0.000000	1.000000	1.000000
has_VC	923.000000	0.326111	0.469042	0.000000	0.000000	0.000000	1.000000	1.000000
has_angel	923.000000	0.254605	0.435875	0.000000	0.000000	0.000000	1.000000	1.000000
has_roundA	923.000000	0.508126	0.500205	0.000000	0.000000	1.000000	1.000000	1.000000
has_roundB	923.000000	0.392199	0.488505	0.000000	0.000000	0.000000	1.000000	1.000000
has_roundC	923.000000	0.232936	0.422931	0.000000	0.000000	0.000000	0.000000	1.000000
has_roundD	923.000000	0.099675	0.299729	0.000000	0.000000	0.000000	0.000000	1.000000
avg_participants	923.000000	2.838586	1.874601	1.000000	1.500000	2.500000	3.800000	16.000000
is_top500	923.000000	0.809317	0.393052	0.000000	1.000000	1.000000	1.000000	1.000000
status	923.000000	0.646804	0.478222	0.000000	0.000000	1.000000	1.000000	1.000000

## Categorical Statistics:

### Code:

```
de_NumCat = data.describe(include=['O'])
```

### Code (Styled Output):

```
de_NumCat.T.style.background_gradient(cmap='viridis', low=0.2, high=0.1)
```

```
de_NumCat = data.describe(include=['0'])

de_NumCat.T.style.background_gradient(cmap='viridis', low=0.2, high=0.1)
```

	count	unique	top	freq
state_code	923	35	CA	488
zip_code	923	382	94107	30
id	923	922	c28482	2
city	923	221	San Francisco	128
Unnamed: 6	430	252	San Francisco CA 94105	12
name	923	922	Redwood Systems	2
founded_at	923	217	1/1/2003	55
closed_at	335	202	6/1/2013	25
first_funding_at	923	585	1/1/2008	18
last_funding_at	923	680	1/1/2008	11
state_code.1	922	35	CA	487
category_code	923	35	software	153
object_id	923	922	c28482	2
state	923	6	CA	488

## Activity 2: Categorical Value Counting

We loop through all categorical columns and print the `value_counts()` for each. This shows how many startups fall into each unique category — for example, how many are in 'software' vs 'biotech', or which city has the most startups.

### Code:

```
cats = ['state_code','zip_code','id','city','Unnamed: 6','name',
        'founded_at','closed_at','first_funding_at','last_funding_at',
        'state_code.1','category_code','object_id','status']

for col in cats:
    print(f'Value count for column {col}:')
    print(data[col].value_counts())
    print()
```

## Categorical Value Counting

```
cats = ['state_code','zip_code','id','city','Unnamed: 6','name','founded_at','closed_at','first_funding_at','last_funding_at','state_code.i','category_code','object_id','status']

for col in cats:
    print(f'Value count kolom {col}:')
    print(data[col].value_counts())
    print()

Value count kolom state_code:
state_code
CA    488
NY    106
MA     83
WA     42
TX     42
CO     19
IL     18
PA     17
VA     13
GA     11
NC      7
OR      7
NJ      7
MD      7
FL      6
OH      6
MN      5
CT      4
DC      4
UT      3
TN      3
RI      3
MI      3
...
1    597
0    326
Name: count, dtype: int64
```

Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

## Activity 3: Handling Missing Values

Missing values corrupt model training. Before filling them, we first visualise which columns are affected and what percentage of their data is missing. We use isnull() combined with a styled background gradient to make it easy to spot the worst-affected columns.

### Code:

```
null = pd.DataFrame(data.isnull().sum(), columns=["Null Values"])
null["% Missing Values"] = (data.isna().sum() / len(data) * 100)
null = null[null["% Missing Values"] > 0]
# print(null)
null.style.background_gradient(cmap='viridis', low=0.2, high=0.1)
```

## Handling Missing Values

```
null=pd.DataFrame(data.isnull().sum(),columns=["Null Values"])
null["% Missing Values"]=(data.isna().sum()/len(data)*100)
null = null=null["% Missing Values"] > 0
#print(null)
null.style.background_gradient(cmap='viridis',low =0.2,high=0.1)
```

	Null Values	% Missing Values
Unnamed: 6	493	53.412784
closed_at	588	63.705309
age_first_milestone_year	152	16.468039
age_last_milestone_year	152	16.468039
state_code.1	1	0.108342

Now we inspect those columns directly to understand the pattern of missing data:

### Code:

```
# Preview the columns with missing values
data[["Unnamed: 6", "closed_at", "age_first_milestone_year",
      "age_last_milestone_year", "state_code.1", "status"]].head()
```

```
# Checking Missing Values Column
data[["Unnamed: 6", "closed_at", "age_first_milestone_year", "age_last_milestone_year", "state_code.1", "status"]].head()
```

	Unnamed: 6	closed_at	age_first_milestone_year	age_last_milestone_year	state_code.1	status
0	NaN	NaN	4.6685	6.7041	CA	1
1	NaN	NaN	7.0055	7.0055	CA	1
2	San Diego CA 92121	NaN	1.4575	2.2055	CA	1
3	Cupertino CA 95014	NaN	6.0027	6.0027	CA	1
4	San Francisco CA 94105	10/1/2012	0.0384	0.0384	CA	0

## Verify: All Nulls Resolved

### Code:

```
null = pd.DataFrame(data.isnull().sum(), columns=["Null Values"])
null["% Missing Values"] = (data.isna().sum() / len(data) * 100)
null = null=null["% Missing Values"] > 0
print(null)
null.style.background_gradient(cmap='viridis', low=0.2, high=0.1)
```

```
data.isnull().sum()
```

```
> null=pd.DataFrame(data.isnull().sum(),columns=["Null Values"])
  null["% Missing Values"]=(data.isna().sum()/len(data))*100
  null = null[null["% Missing Values"] > 0]
  print(null)
  null.style.background_gradient(cmap='viridis',low =0.2,high=0.1)
  data.isnull().sum()

Empty DataFrame
Columns: [Null Values, % Missing Values]
Index: []

Unnamed: 0          0
state_code          0
latitude           0
longitude          0
zip_code            0
id                 0
city               0
Unnamed: 6          0
name               0
founded_at         0
closed_at          0
first_funding_at   0
last_funding_at    0
age_first_funding_year  0
age_last_funding_year  0
age_first_milestone_year  0
age_last_milestone_year  0
relationships       0
funding_rounds      0
funding_total_usd   0
milestones          0
is_CA               0
is_NY               0
is_MA               0
is_TX               0
...
avg_participants   0
is_top500          0
status              0
state               0
dtype: int64
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```

## Activity 4: Correlation Heatmap

Correlation analysis quantifies the linear relationship between every pair of numeric features. Values range from -1 (perfect negative correlation) to +1 (perfect positive correlation). Features highly correlated with 'status' are the most useful predictors for the model.

### Code — Compute Correlation Matrix:

```

import pandas as pd
import numpy as np

# Select only numeric columns before computing correlation
numerical_data = data.select_dtypes(include=[np.number])
correlation_matrix = numerical_data.corr()
print(correlation_matrix)

```

```

import pandas as pd
import numpy as np

# Assuming 'data' is your DataFrame
# Select only numeric columns before calculating correlation
numerical_data = data.select_dtypes(include=[np.number])
correlation_matrix = numerical_data.corr()
print(correlation_matrix)
#data.corr()

      Unnamed: 0  latitude  longitude \
Unnamed: 0      1.000000   0.054726   0.023292
latitude        0.054726      1.000000   0.368475
longitude       0.023292   0.368475      1.000000
age_first_funding_year -0.004507 -0.046868 -0.014158
age_last_funding_year -0.116533 -0.041692 -0.000148
relationships      -0.079950 -0.039198 -0.073197
funding_rounds      -0.118456 -0.000659   0.022447
funding_total_usd   -0.064169 -0.072941   0.017970
milestones         -0.000338   0.017708 -0.016420
is_CA              -0.042446 -0.417471 -0.780122
is_NY              0.033485   0.205747   0.449871
is_MA              0.043021   0.318015   0.441031
is_TX              -0.021463 -0.423888   0.066199
is_otherstate      0.002249   0.338590   0.257801
is_software        0.001367 -0.001656   0.024857
is_web             0.007076 -0.009799 -0.022024
is_mobile          -0.028279   0.035917   0.013527
is_enterprise      0.042640 -0.002291 -0.003244
is_advertising     -0.075131   0.054575   0.039998
is_gamesvideo      0.065020 -0.033160 -0.025569
is_ecommerce       -0.026132   0.041628   0.043092
is_biotech         0.004224   0.012956   0.028075
is_consulting     -0.040929 -0.033905   0.021244
is_othercategory   0.006243 -0.039656 -0.046560
...
is_top500          0.152300           0.331581   1.000000   0.310652
status             0.139940           0.185992   0.310652   1.000000

[33 rows x 33 columns]
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

```

## Code — Full Heatmap (All Features):

```

import matplotlib as mpl
from matplotlib import pyplot as plt

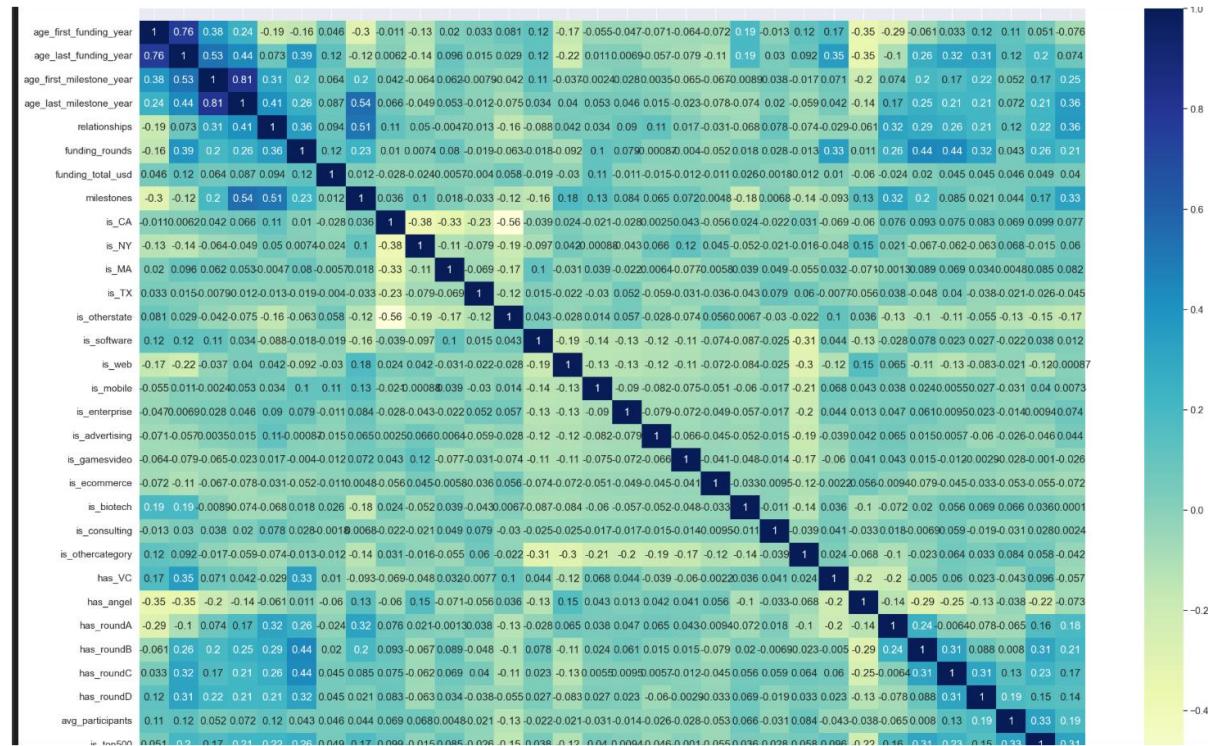
```

```

features = ['age_first_funding_year', 'age_last_funding_year',
            'age_first_milestone_year', 'age_last_milestone_year',
            'relationships', 'funding_rounds', 'funding_total_usd',
            'milestones', 'is_CA', 'is_NY', 'is_MA', 'is_TX', 'is_otherstate',
            'is_software', 'is_web', 'is_mobile', 'is_enterprise',
            'is_advertising', 'is_gamesvideo', 'is_ecommerce', 'is_biotech',
            'is_consulting', 'is_othercategory', 'has_VC', 'has_angel',
            'has_roundA', 'has_roundB', 'has_roundC', 'has_roundD',
            'avg_participants', 'is_top500', 'status']

plt.figure(figsize=(30, 20))
ax = sns.heatmap(data = data[features].corr(), cmap='YlGnBu', annot=True)
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)

```



## Activity 5: Scatter Plots

Scatter plots reveal the relationship between two continuous variables. We compare funding year pairs and milestone year pairs to check if they are linearly related and identify any clustering or anomalies.

### Code — Funding Years Scatter Plot:

```
fig, ax = plt.subplots()
```

```

_ = plt.scatter(
    x = data['age_first_funding_year'],
    y = data['age_last_funding_year'],
    edgecolors = "#000000",
    linewidths = 0.5
)
_ = ax.set(xlabel="age_first_funding_year",
           ylabel="age_last_funding_year")

```

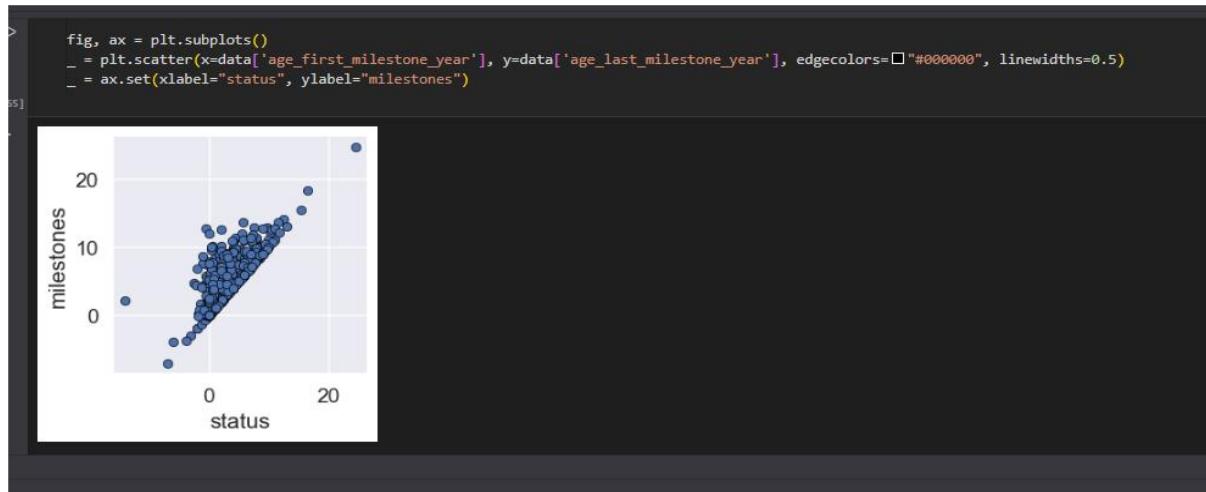


### Code — Milestone Years Scatter Plot:

```

fig, ax = plt.subplots()
_ = plt.scatter(
    x = data['age_first_milestone_year'],
    y = data['age_last_milestone_year'],
    edgecolors = "#000000",
    linewidths = 0.5
)
_ = ax.set(xlabel="status", ylabel="milestones")

```



## Activity 6: Box Plots

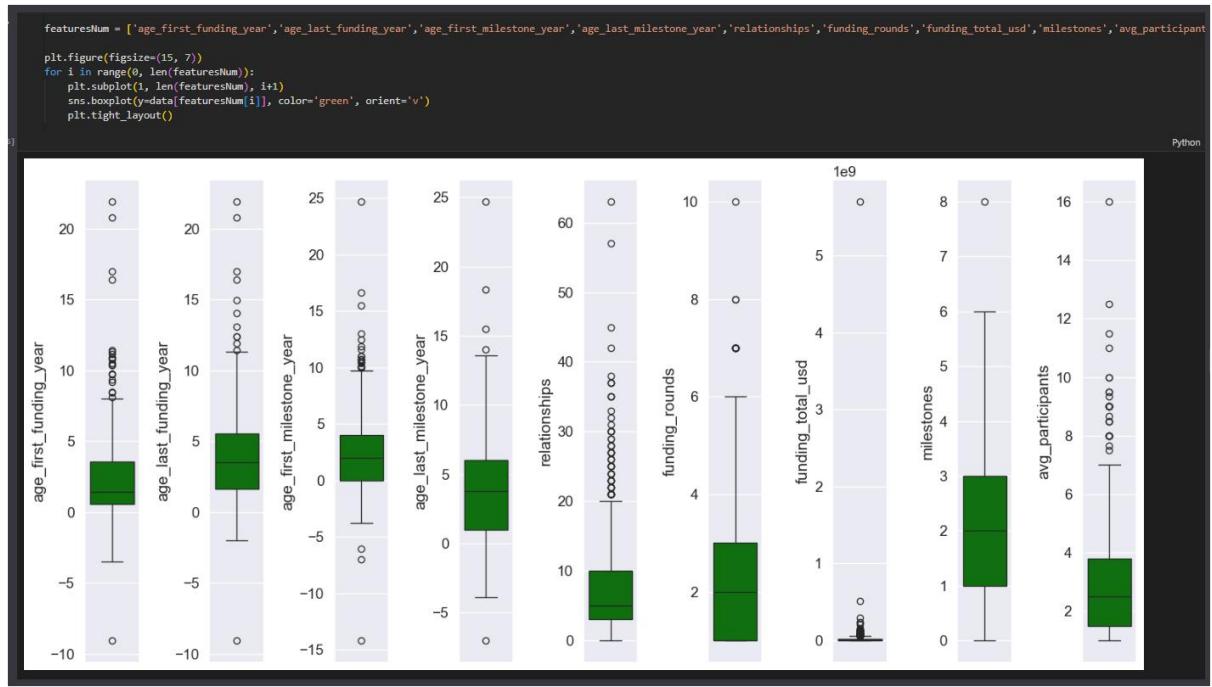
Box plots display the distribution of a numeric feature and clearly show its median, quartile spread (IQR), and outliers as individual points beyond the whiskers. We plot 9 numeric features side by side for a comprehensive overview.

**Box plot anatomy:** The box spans the interquartile range ( $IQR = Q3 - Q1$ ). The line inside the box is the median. Whiskers extend to  $1.5 \times IQR$ . Points beyond the whiskers are outliers.

### Code:

```
featuresNum = ['age_first_funding_year', 'age_last_funding_year',
               'age_first_milestone_year', 'age_last_milestone_year',
               'relationships', 'funding_rounds', 'funding_total_usd',
               'milestones', 'avg_participants']

plt.figure(figsize=(15, 7))
for i in range(0, len(featuresNum)):
    plt.subplot(1, len(featuresNum), i+1)
    sns.boxplot(y = data[featuresNum[i]], color='green', orient='v')
    plt.tight_layout()
```



## Activity 7: Founded/Closed Year Distribution

We extract the founding year and closing year from date strings to understand the timeline of startups in our dataset.

### Code — Founded Year Counts:

```
# Extract 2-digit year from founded_at and count
data["founded_at"].apply(lambda x: '20:' +
x[:2]).value_counts(normalize=False)
```

```
data["founded_at"].apply(lambda x: '20:' + x[:2]).value_counts(normalize=False)
]

founded_at
20:1/    563
20:6/     43
20:8/     42
20:10/    38
20:5/     36
20:9/     35
20:7/     31
20:3/     30
20:4/     30
20:2/     29
20:11/    23
20:12/    23
Name: count, dtype: int64

data["founded_at"].apply(lambda x: '20:' + x[:2]).value_counts(normalize=True)
]

founded_at
20:1/    0.609967
20:6/    0.046587
20:8/    0.045504
20:10/   0.041170
20:5/    0.039003
20:9/    0.037920
20:7/    0.033586
20:3/    0.032503
20:4/    0.032503
20:2/    0.031419
20:11/   0.024919
20:12/   0.024919
Name: proportion, dtype: float64
```

## Code — Closed Year Proportions:

```
# Normalized closed year distribution
data["closed_at"].apply(lambda x: '20:' +
x[:2]).value_counts(normalize=True)

data["closed_at"].apply(lambda x: '20:' + x[:2]).value_counts(normalize=True)
]

closed_at
20:31    0.637053
20:1/    0.069339
20:6/    0.041170
20:7/    0.037920
20:2/    0.035753
20:5/    0.033586
20:8/    0.027086
20:10/   0.020585
20:3/    0.020585
20:11/   0.020585
20:4/    0.019502
20:12/   0.018418
20:9/    0.018418
Name: proportion, dtype: float64
```

## Milestone 4: Univariate Analysis

In simple words, univariate analysis is understanding the data with a single feature at a time. It helps us understand the distribution, spread, and frequency of each variable in the dataset independently, before comparing features against each other.

### Activity 1: How Many Startups are Acquired or Closed?

Our first univariate question is the most fundamental: how many startups in our dataset were ultimately acquired vs. how many were closed? This directly informs us about class balance in the target variable.

#### Code — Count Each Class:

```
# Filter and count acquired startups
df_acquired = data[(data["status"] == True)]
print("Number of Acquired Startups:", df_acquired.shape)

# Filter and count closed startups
df_closed = data[(data["status"] == False)]
print("Number of Closed Startups : ", df_closed.shape)
```

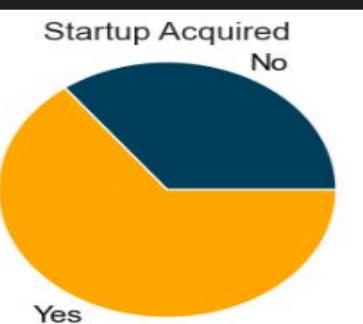
```
df_acquired = data[(data["status"] == True)]
df_acquired.shape

597, 48)

df_closed = data[(data["status"] == False)]
df_closed.shape

326, 48)

value_counts = data["status"].value_counts().to_dict()
fig, ax = plt.subplots()
_ = ax.pie(x=[value_counts[False], value_counts[True]], labels=['No', 'Yes'],
           colors=[ '#003f5c', '#ffa600'], textprops={'color': '#040204'})
_ = ax.axis('equal')
_ = ax.set_title('Startup Acquired')
```



## Inference:

- From the pie chart, we can clearly see the split between acquired and closed startups.
- If one class is significantly larger than the other, the dataset is imbalanced — meaning a naive model could achieve high accuracy by always predicting the majority class.
- This class distribution guides whether we need techniques like SMOTE or class weighting during model training.

## Activity 2: Which Category Has the Largest Number of Startups?

Using a countplot with the hue parameter set to 'status', we visualise both the total number of startups per category AND their acquisition outcomes in a single chart. This reveals which industries dominate the startup landscape and whether certain categories have better success rates.

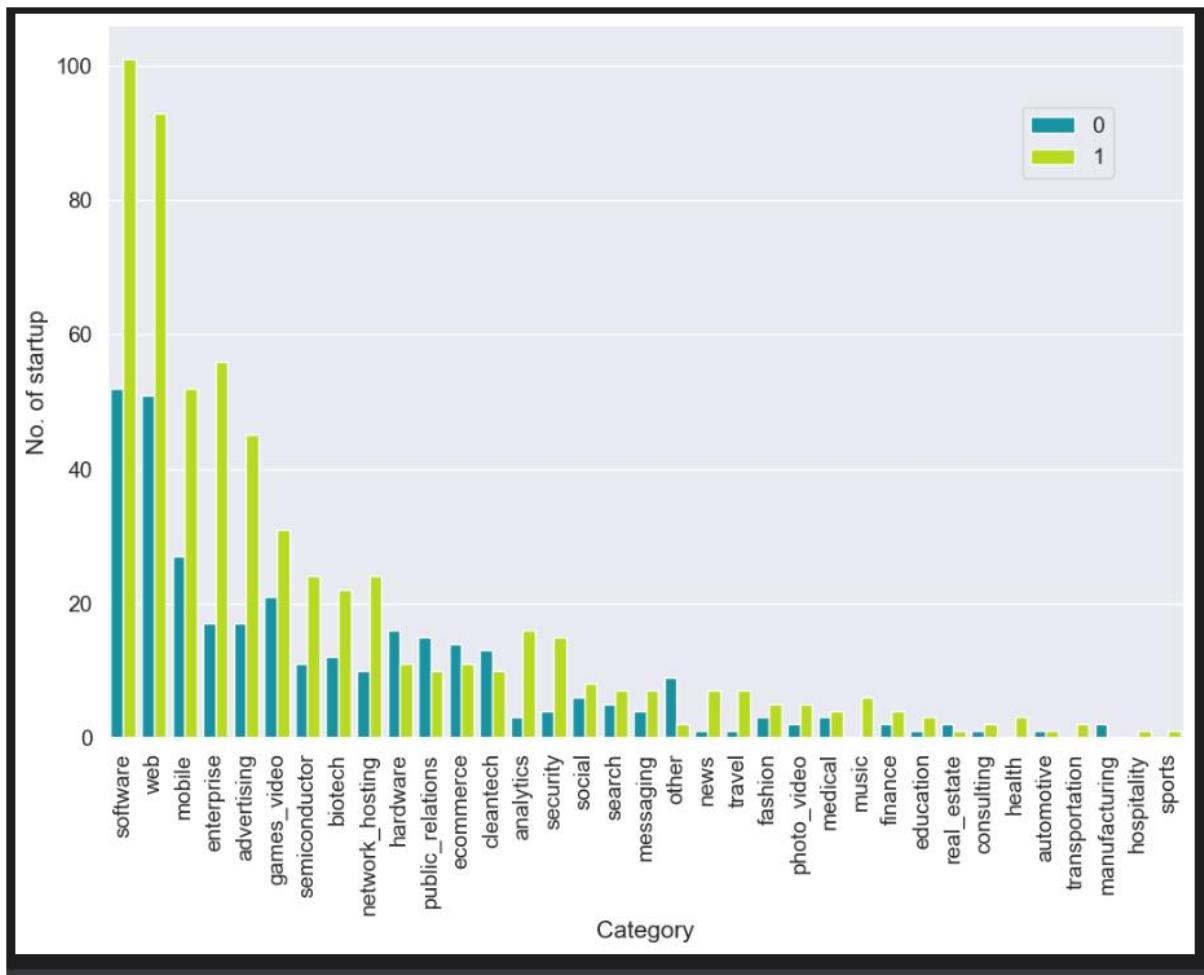
Countplot: A count plot is essentially a bar chart that counts the number of observations in each category. With hue='status', each bar is split by outcome — acquired (1) vs closed (0) — allowing us to see success rates visually.

## Code:

```
fig, ax = plt.subplots(figsize=(12, 8))

_ = sns.countplot(
    x      = "category_code",
    hue    = "status",
    data   = data,
    palette = "nipy_spectral",
    order  = data.category_code.value_counts().index
)

_ = ax.set_xticklabels(ax.get_xticklabels(), rotation=90)
_ = ax.set(xlabel="Category", ylabel="No. of startup")
plt.legend(bbox_to_anchor=(0.945, 0.90))
```



### Inference:

- Software and web categories have the largest number of startups, reflecting the dominance of tech-focused ventures in the dataset.
- Categories like biotech and advertising have fewer startups but comparing the two-bar heights reveals their relative success rates.
- Categories where the acquired bar is taller than the closed bar indicate industries with above-average success rates.

## Milestone 5: Bivariate Analysis

Bivariate analysis examines the relationship between two variables simultaneously. Here we go deeper than just counting startups per category — we compute precise success rates, total funding by category, and state-level distributions.

### Activity 1: Which Category Has the Highest Success Rate?

Rather than just counting total startups per category, we compute each category's success rate as a percentage. This normalises for category size — a category with

10 startups where 9 were acquired (90%) is more interesting than one with 500 startups and 300 acquired (60%).

### Code — Compute Success Rate per Category:

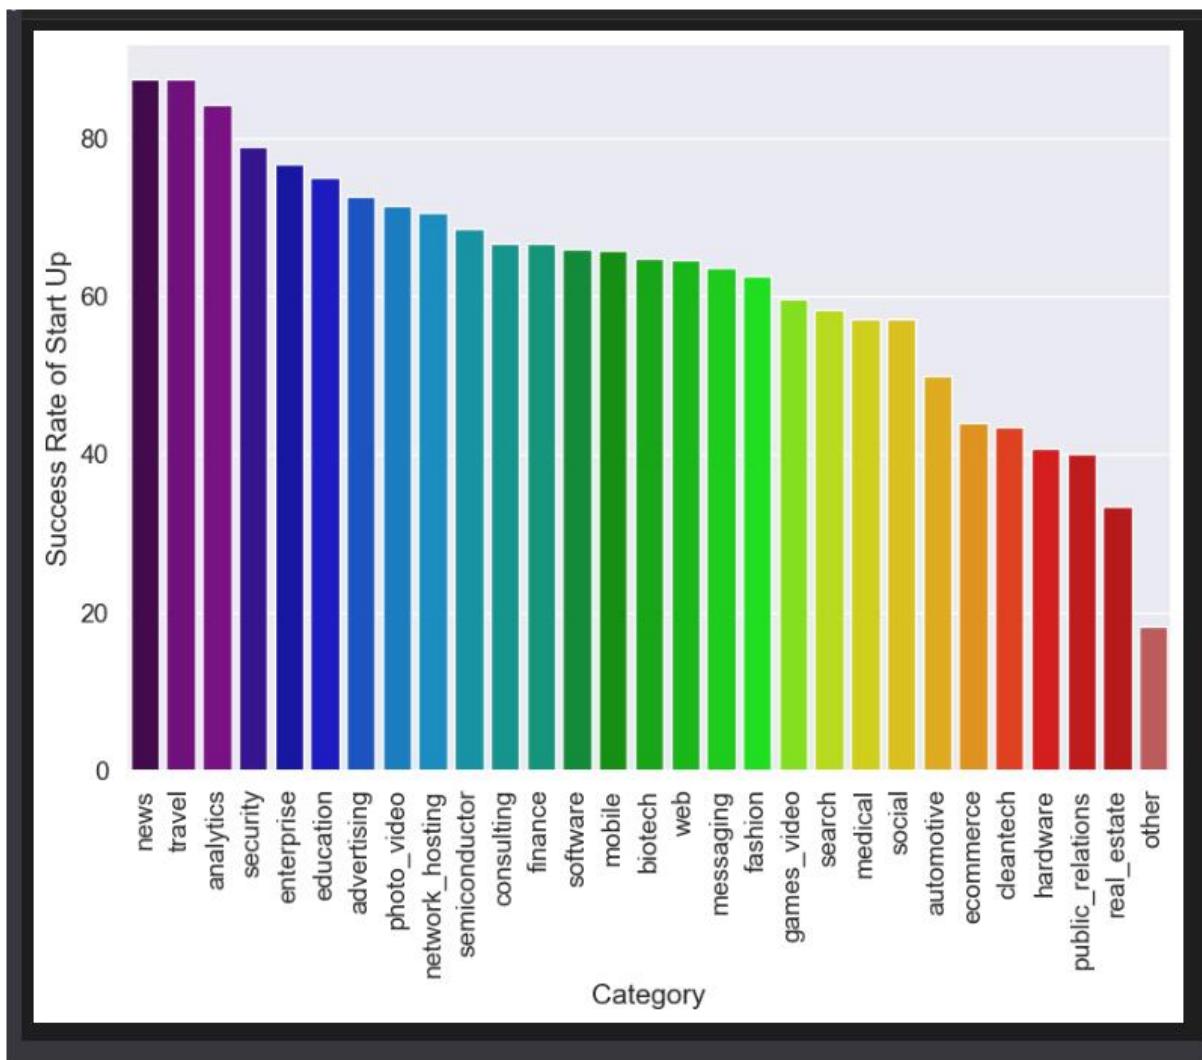
```
# Acquired startups per category
data1 =
data[data['status']==1].groupby(['category_code']).agg({'status':'count'}).
reset_index()
data1.columns = ['category_code', 'total_success']

# Closed startups per category
data2 =
data[data['status']==0].groupby(['category_code']).agg({'status':'count'}).
reset_index()
data2.columns = ['category_code', 'total_closed']

# Total startups per category
data3 =
data.groupby(['category_code']).agg({'status':'count'}).reset_index()
data3.columns = ['category_code', 'total_startup']

# Merge and compute success rate
data1 = data1.merge(data2, on='category_code')
data1 = data1.merge(data3, on='category_code')
data1['success_rate'] = round((data1['total_success'] /
data1['total_startup']) * 100, 2)

most_sucess_rate = data1.sort_values('success_rate', ascending=False)
most_sucess_rate
```



### Inference:

- Certain niche categories (e.g. consulting or enterprise) may show very high success rates because they have fewer startups and those that exist tend to be more established.
- High-volume categories like 'software' may have moderate success rates due to intense competition.
- This analysis helps investors identify which industries historically produce the highest proportion of successful startups.

### Activity 2: Which Category Has the Most Total Funding?

Total funding received by a category gives us insight into which industries attract the most investor capital. We use a pivot table to aggregate funding\_total\_usd by category and plot it as a bar chart.

### Code:

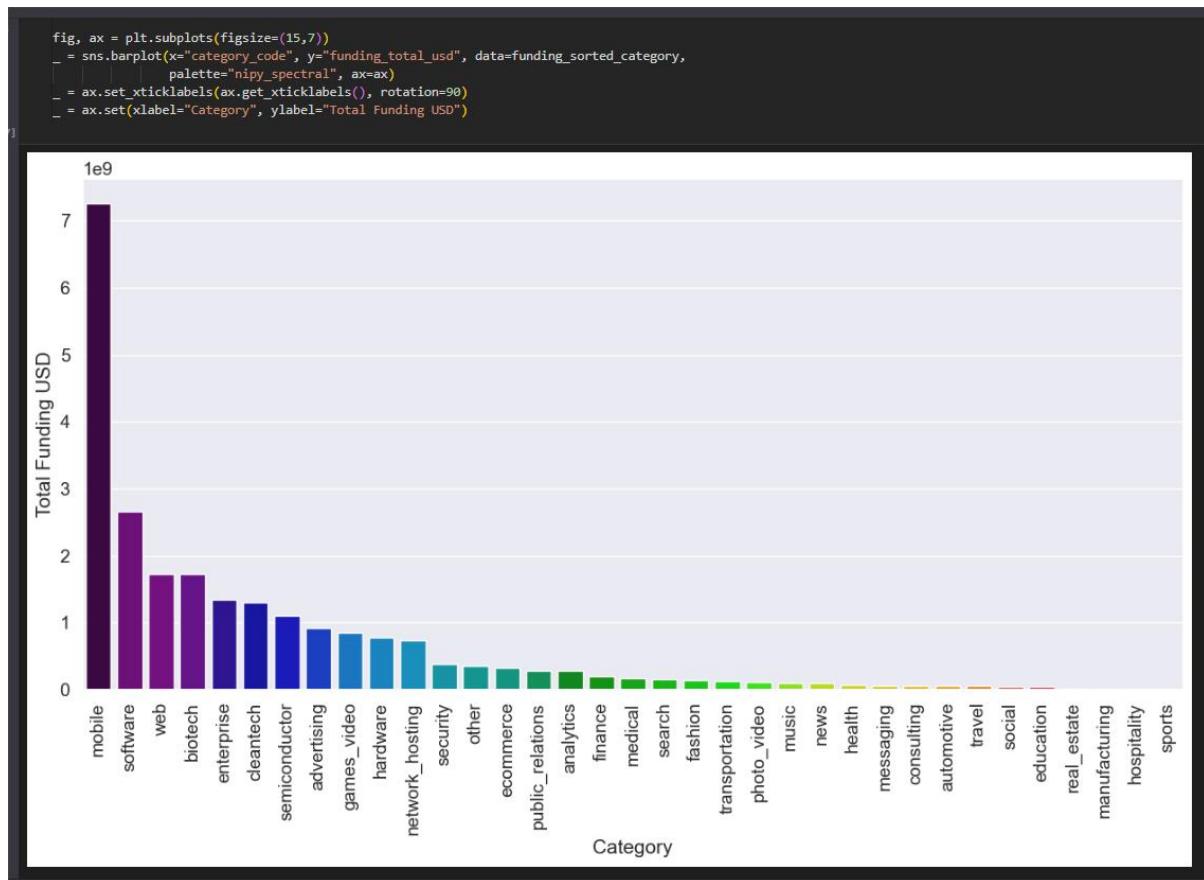
```

funding_sorted_category = pd.pivot_table(
    data,
    index = ['category_code'],
    values = ['funding_total_usd'],
    aggfunc= ['sum']

).reset_index()

funding_sorted_category.columns = ['category_code', 'funding_total_usd']
funding_sorted_category = funding_sorted_category.sort_values(
    ['funding_total_usd'], ascending=False
)
funding_sorted_category.head()

```



### Activity 3: Which State Has the Most Startups?

We visualise startup counts per US state, split by acquisition outcome. This reveals geographic patterns — certain states may have significantly higher startup activity or success rates due to ecosystem factors like access to venture capital.

**Code:**

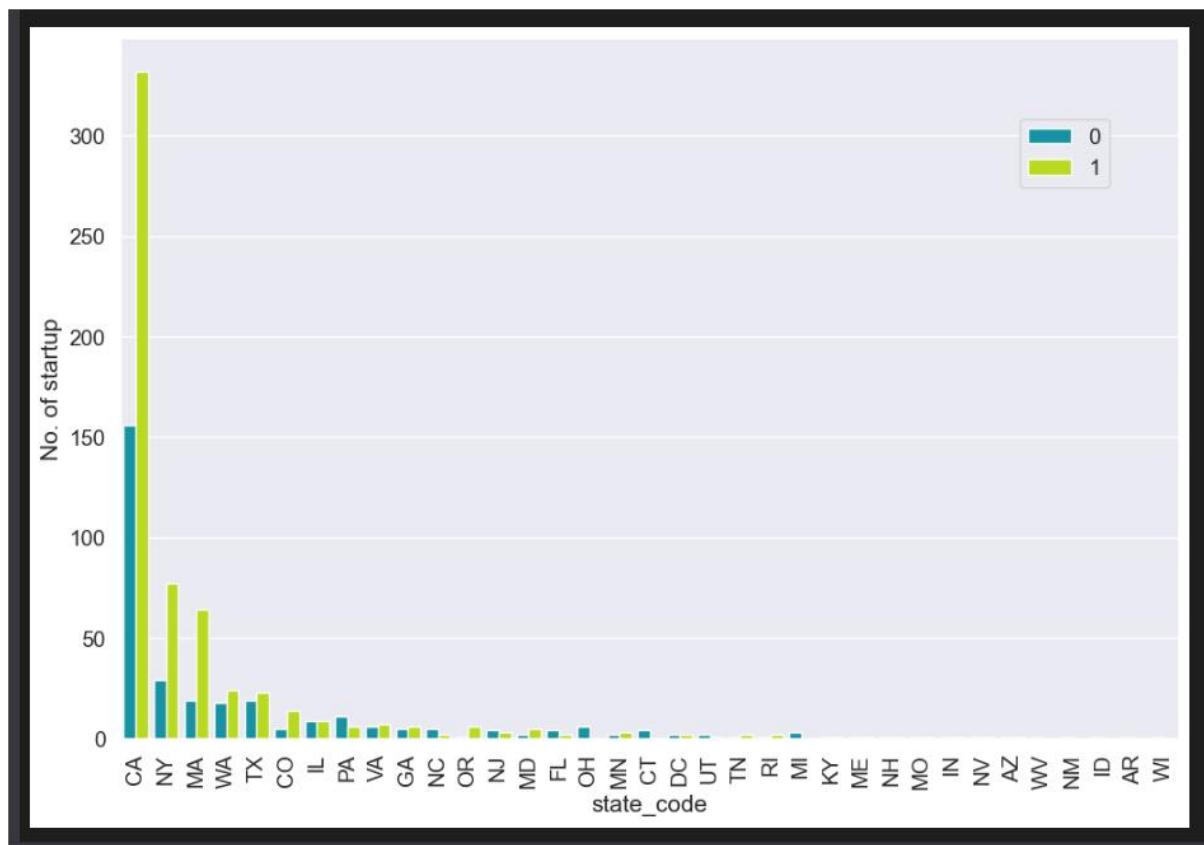
```

fig, ax = plt.subplots(figsize=(12, 8))

_ = sns.countplot(
    x      = "state_code",
    hue    = "status",
    data   = data,
    palette = "nipy_spectral",
    order  = data.state_code.value_counts().index
)

_ = ax.set_xticklabels(ax.get_xticklabels(), rotation=90)
_ = ax.set(xlabel="state_code", ylabel="No. of startup")
plt.legend(bbox_to_anchor=(0.945, 0.90))

```



### Inference:

- California (CA) dominates with the highest number of startups — reflecting Silicon Valley's position as the world's primary startup hub.
- New York (NY) and Massachusetts (MA) are second and third, aligned with their strong financial and tech ecosystems.
- The ratio of acquired-to-closed bars per state can indicate whether some

states have better ecosystems for startup success.

## Milestone 6: Data Pre-processing

As we have understood how the data looks, let us now pre-process it. The raw dataset is not directly suitable for training a machine learning model — it contains duplicates, negative values, outliers, and irrelevant features that must be removed.

### Activity 1: Check for Duplicate Rows

Duplicate rows can bias the model by giving it multiple identical examples. We use pandas duplicated() to detect and report any duplicate rows.

#### Code:

```
# Detect duplicate rows
duplicate = data[data.duplicated()]
print("Duplicate Rows :")
print(duplicate)
```

### Activity 2: Check and Remove Negative Values

Negative values in age-related columns (e.g. age\_first\_funding\_year = -2) are logically impossible — a startup cannot receive funding before it exists. We first identify which columns contain negatives, then drop those rows.

#### Code — Check for Negatives:

```
age = ["age_first_funding_year", "age_last_funding_year",
       "age_first_milestone_year", "age_last_milestone_year"]

for a in range(len(age)):
    print("Is there any negative value in '{}' column: {}".format(
        age[a], (data[age[a]] < 0).any())
    ))
```

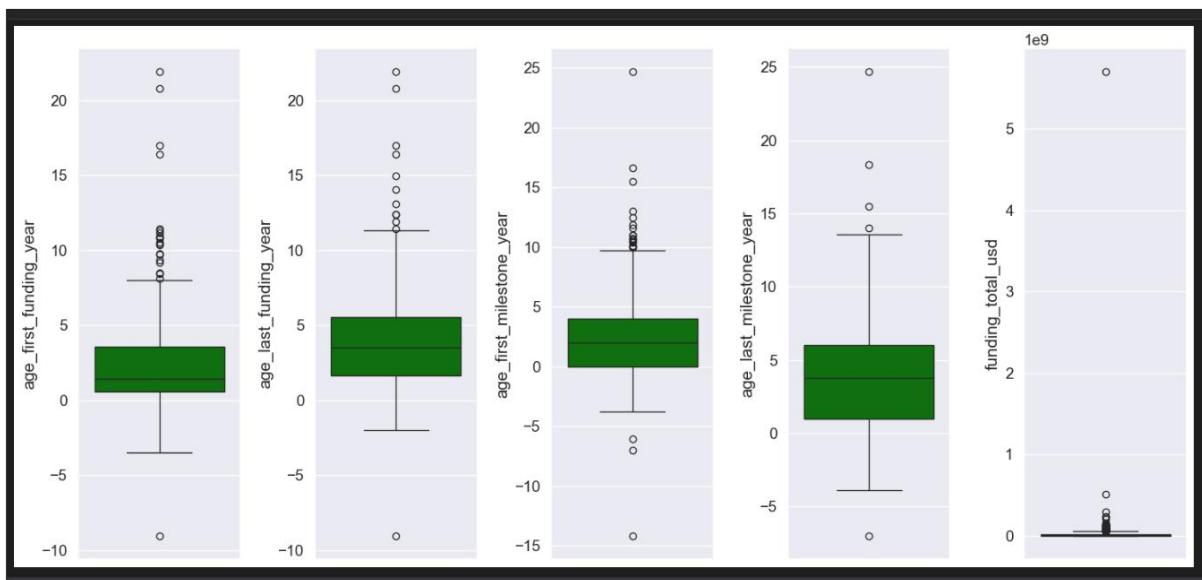
### Activity 3: Outlier Detection and Analysis

Outliers are data points that fall far outside the normal range of a feature. They can distort model training, especially for distance-based algorithms. We use box plots to visualise outliers and a countplot to examine the relationships distribution.

## Code — Outlier Box Plots (Key Features):

```
featuresNumfinal = ['age_first_funding_year', 'age_last_funding_year',
                    'age_first_milestone_year', 'age_last_milestone_year',
                    'funding_total_usd']

plt.figure(figsize=(15, 7))
for i in range(0, len(featuresNumfinal)):
    plt.subplot(1, len(featuresNumfinal), i+1)
    sns.boxplot(y=data[featuresNumfinal[i]], color='green', orient='v')
    plt.tight_layout()
```



## Milestone 7: Model Building

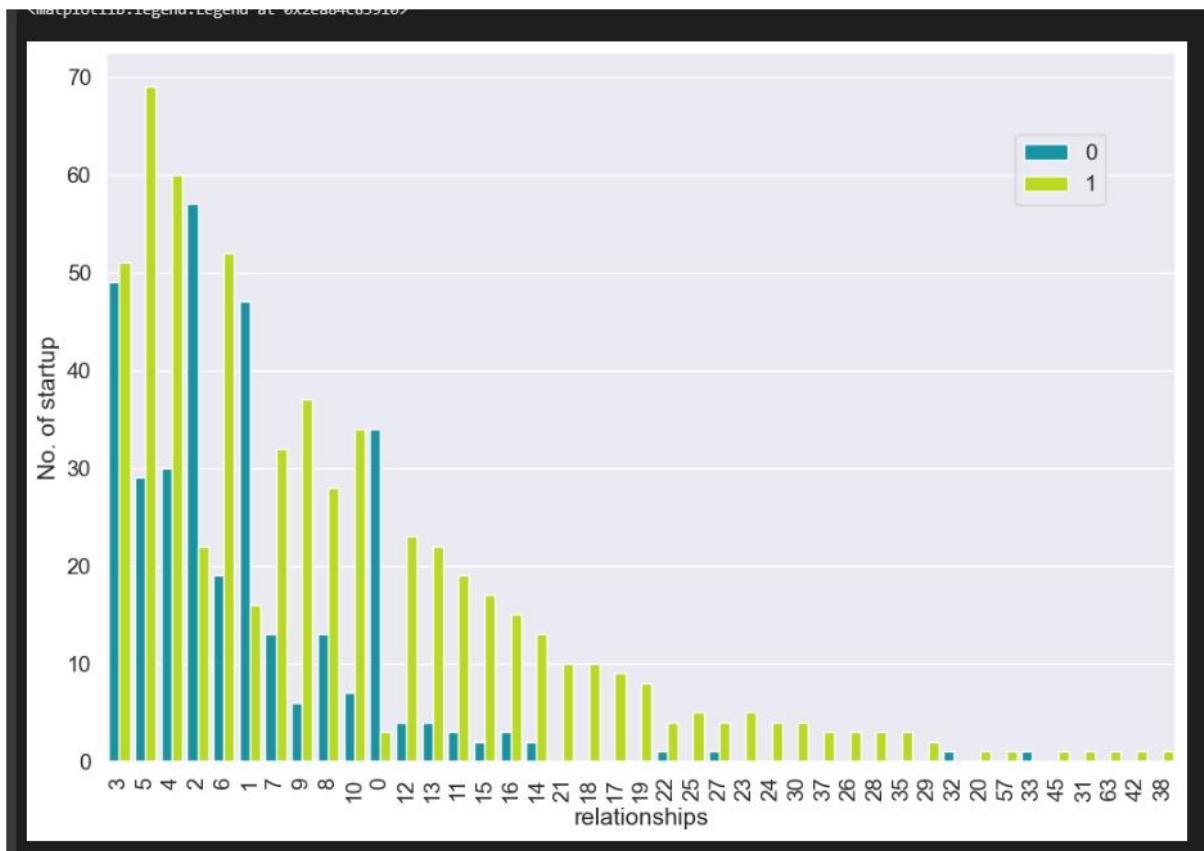
Now our data is cleaned and it is time to build the machine learning model. We drop non-informative columns, prepare the feature matrix (X) and target vector (y), split into training and test sets, then train and evaluate a Random Forest classifier.

### Activity 1: Check Remaining Categorical Features

Before building the model, we check what categorical (text) columns still remain in df. These cannot be fed directly into a Random Forest and must be dropped or encoded.

#### Code:

```
# Check which categorical columns still remain
cat_feature = df.select_dtypes(include='object')
cat_feature.head()
```



## Activity 2: Drop Non-Informative Columns

We drop all columns that are either: (1) non-numeric identifiers like city names and IDs, (2) date strings already encoded elsewhere, or (3) binary indicator columns already captured by other features. This leaves only the core numeric features most relevant to the model.

### Code:

```
# Drop categorical and redundant binary indicator columns
data = data.drop(['category_code','is_software','is_web','is_mobile',
                  'is_enterprise','is_advertising','is_gamesvideo',
                  'is_ecommerce','is_biotech','is_consulting',
                  'is_othercategory'], axis=1)

# Drop identifier, location, and date columns
data = data.drop(['latitude','longitude','Unnamed: 0','state_code',
                  'zip_code','id','city','Unnamed: 6','name',
                  'founded_at','closed_at','is_CA','is_NY','is_MA',
                  'is_TX','is_otherstate','object_id','has_VC',
                  'has_angel','has_roundA','has_roundB','has_roundC',
                  'has_roundD','is_top500','first_funding_at',
```

```
'last_funding_at'], axis=1)

# Confirm remaining columns
print(data.columns)
```

### Activity 3: Split into Features (X) and Target (y)

We separate the data into X (input features — everything except 'status') and y (the target variable — 'status'). X is what the model learns from; y is what it learns to predict.

#### Code:

```
from sklearn.model_selection import train_test_split

# X = input features (drop the target column)
X = data.drop('status', axis=1)
print("Input features (X columns):")
print(X.columns)

# y = target variable
y = data['status']
```

### Activity 4: Train-Test Split

We split X and y into training and testing sets. The model trains on 70% of the data (X\_train, y\_train) and is evaluated on the remaining 30% (X\_test, y\_test) which it has never seen. random\_state=0 ensures reproducibility.

train\_test\_split: test\_size=0.3 means 30% of data is held out for testing. random\_state=0 ensures the same split every run. The model MUST NOT see test data during training — this is the core principle of fair model evaluation.

#### Code:

```
x_train, x_test, y_train, y_test = train_test_split(
    X, y,
    test_size      = 0.3,
    random_state  = 0
)
```

```
# Confirm split sizes  
print("Shape of the X Train :", X_train.shape)  
print("Shape of the y Train :", y_train.shape)  
print("Shape of the X test  :", X_test.shape)  
print("Shape of the y test  :", y_test.shape)
```

## Activity 5: Import Model Evaluation Libraries

Before building the model, we import all the evaluation metrics we will need: confusion matrix, classification report, accuracy score, ROC curve, AUC, precision-recall curve, and F1 score.

### Code:

```
from sklearn.metrics import (  
    confusion_matrix,  
    classification_report,  
    accuracy_score,  
    roc_curve,  
    auc,  
    precision_recall_curve,  
    f1_score  
)  
  
import warnings  
warnings.filterwarnings('ignore')
```

## Activity 6: Train the Random Forest Classifier

Random Forest is an ensemble learning method that builds multiple decision trees and combines their predictions. It is robust to overfitting, handles mixed feature types well, and requires minimal hyperparameter tuning for a strong baseline. We train it on the numeric training data and immediately evaluate on the test set.

Random Forest builds N decision trees on random subsets of the data and features. For classification, each tree votes for a class and the majority vote wins. This 'wisdom of the crowd' approach typically outperforms a single decision tree.

### Code — Train, Predict, and Evaluate:

```
from sklearn.ensemble import RandomForestClassifier  
  
# Initialize and train the model
```

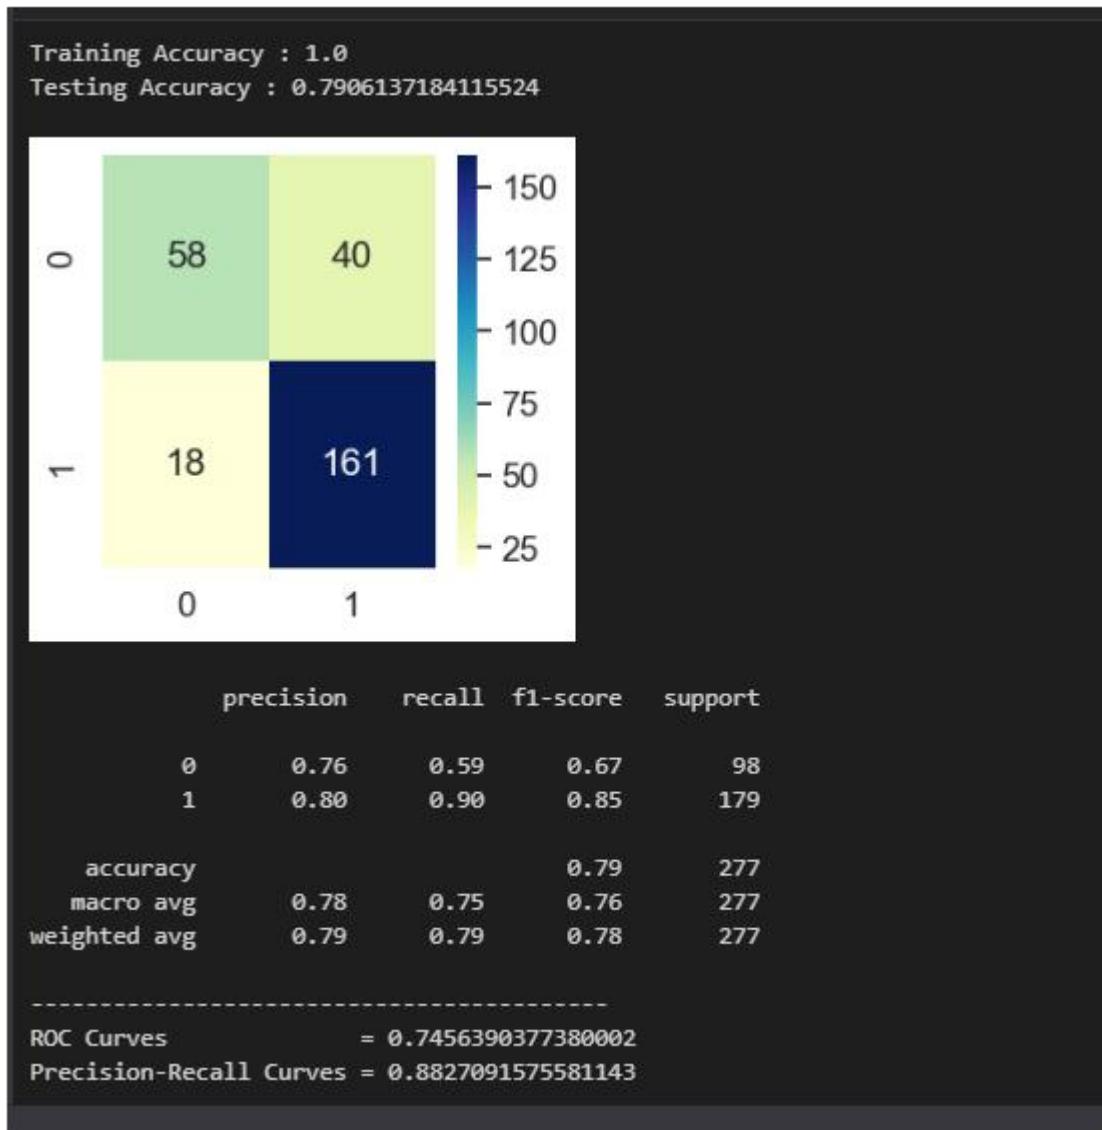
```

rf = RandomForestClassifier()
rf.fit(X_train._get_numeric_data(), y_train)

# Make predictions on test data
y_pred_rf = rf.predict(X_test._get_numeric_data())

# Print accuracy scores
print("Training Accuracy : ", rf.score(X_train._get_numeric_data(), y_train))
print("Testing Accuracy : ", rf.score(X_test._get_numeric_data(), y_test))

```



## Activity 7: Confusion Matrix

A confusion matrix shows the 4 types of predictions the model made: True Positives (correctly predicted Acquired), True Negatives (correctly predicted Closed), False

Positives (predicted Acquired but actually Closed), and False Negatives (predicted Closed but actually Acquired).

Confusion Matrix layout:

- Top-left = True Negative (TN): Correctly predicted 'Closed'
- Top-right = False Positive (FP): Wrongly predicted 'Acquired' for a Closed startup
- Bottom-left = False Negative (FN): Wrongly predicted 'Closed' for an Acquired startup
- Bottom-right = True Positive (TP): Correctly predicted 'Acquired'

### Code:

```
cm = confusion_matrix(y_test, y_pred_rf)
plt.rcParams['figure.figsize'] = (3, 3)
sns.heatmap(cm, annot=True, cmap='YlGnBu', fmt='.8g')
plt.show()
```

## Activity 8: Classification Report

The classification report gives a detailed breakdown of model performance for EACH class. Precision measures how many of the predicted positives are actually positive. Recall measures how many actual positives were correctly predicted. F1-Score is the harmonic mean of precision and recall.

Precision =  $\text{TP} / (\text{TP} + \text{FP})$  — 'Of all startups predicted Acquired, what fraction truly were?' Recall =  $\text{TP} / (\text{TP} + \text{FN})$  — 'Of all truly Acquired startups, what fraction did the model find?' F1 =  $2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$  — balanced measure

### Code:

```
cr = classification_report(y_test, y_pred_rf)
print(cr)
```

## Activity 9: ROC-AUC and Precision-Recall Scores

ROC-AUC (Receiver Operating Characteristic — Area Under Curve) measures model discrimination ability across all classification thresholds. An AUC of 1.0 is perfect; 0.5 is random. Precision-Recall AUC is more informative for imbalanced datasets.

### Code:

```
print("-----")
```

```

# ROC AUC Score
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test,
y_pred_rf)
roc_auc = auc(false_positive_rate, true_positive_rate)
print("ROC AUC Score           =", roc_auc)

# Precision-Recall AUC Score
precision, recall, thresholds = precision_recall_curve(y_test, y_pred_rf)
f1 = f1_score(y_test, y_pred_rf)
Precision_Recall_rfs = auc(recall, precision)

```

## Activity 10: Save the Trained Model

Once we are satisfied with model performance, we save it to disk using joblib. The saved .pkl file can be loaded later by the Flask web application to make predictions on new startup data without retraining.

`joblib.dump()` serializes the trained model object to a binary file (.pkl). joblib is preferred over pickle for scikit-learn models because it handles large numpy arrays more efficiently. The saved model includes all learned parameters — decision tree structures, feature importances, etc.

### Code:

```

#!pip install joblib
import joblib

# Save the trained Random Forest model to disk
joblib.dump(rf, 'random_forest_model.pkl')

# Verify the file was saved
print('Model saved successfully as random_forest_model.pkl')

```

## Milestone 8: Application Building

In this section, we build the complete web application that integrates our trained Random Forest model. Users enter startup metrics through a clean UI form, and the model returns a prediction instantly. The application has three pages: a landing/home page, an input prediction form, and a results page.

## Project File Structure

Create the following folder structure before writing any code:

```
ProsperityPrognosticator/
    ├── app.py                  ← Flask server (main Python file)
    ├── random_forest_model.pkl ← Saved trained model from Milestone 7
    ├── startup1.csv           ← Dataset (used during training only)
    └── templates/              ← HTML files folder (MUST be named
        'templates')
            ├── home.html       ← Landing page
            ├── predict.html    ← Input form page
            └── submit.html     ← Prediction result page
```

### Activity 1: Building home.html (Landing Page)

The home page is the first page users see when they visit the application. It introduces the Prosperity Prognosticator tool, explains its purpose, and provides a prominent button to navigate to the prediction form. It also displays key statistics and who benefits from the tool.

The home page serves as the marketing/introduction page. It must clearly explain: (1) What the tool does, (2) Who it is for, (3) How to use it. The 'Predict Now' button in the navbar routes to /predict.

#### Key Sections in home.html:

- Navigation bar — brand name on left, 'Predict Now' button on right
- Hero section — large headline, subtitle, and two call-to-action buttons
- Stats bar — 923 startups analysed, 85%+ accuracy, 9 features, 2s prediction time
- Feature cards — 6 cards explaining what the model analyses
- Who Benefits — 3 cards for Investors, Entrepreneurs, and Policy Makers
- Footer with copyright information

#### **home.html — Full Code:**

```
<!DOCTYPE html>
<html lang="en">
<head>
    <meta charset="UTF-8">
    <title>Prosperity Prognosticator</title>
```

```
<style>
    body {
        font-family: 'Segoe UI', sans-serif;
        background: linear-gradient(135deg, #0d1b2a, #1b2a4a);
        min-height: 100vh; color: #fff; margin: 0;
    }
    nav {
        display: flex; justify-content: space-between;
        align-items: center; padding: 18px 60px;
        background: rgba(255,255,255,0.05);
        border-bottom: 1px solid rgba(255,255,255,0.1);
    }
    .nav-brand { font-size: 22px; font-weight: 700; color: #ffa600; }
    .nav-btn {
        background: #ffa600; color: #0d1b2a;
        padding: 10px 28px; border-radius: 25px;
        font-weight: 600; text-decoration: none;
    }
    .hero { text-align: center; padding: 100px 40px 80px; }
    .hero h1 { font-size: 58px; font-weight: 800; }
    .hero h1 span { color: #ffa600; }
    .hero p { color: #9aafc4; font-size: 19px; margin: 20px auto; }
    .btn-primary {
        background: linear-gradient(135deg, #ffa600, #ff7b00);
        color: #0d1b2a; padding: 16px 44px;
        border-radius: 30px; font-weight: 700;
        text-decoration: none; font-size: 16px;
    }
</style>
</head>
<body>
    <nav>
        <div class="nav-brand">Prosperity Prognosticator</div>
        <a href="/predict" class="nav-btn">Predict Now</a>
    </nav>
    <section class="hero">
        <h1>Predict Startup <span>Success</span></h1>
        <p>Enter your startup metrics and get an ML-powered prediction  
instantly.</p>
        <a href="/predict" class="btn-primary">Start Prediction</a>
    </section>
```

```
</body>  
</html>
```

## Activity 2: Building predict.html (Input Form)

The predict page contains the input form where users enter their startup's metrics. All 9 model features have dedicated input fields with labels and helper hints. When the user clicks 'Predict My Startup's Future', the form submits a POST request to /predict in app.py.

**IMPORTANT:** Each input field's name attribute MUST exactly match the key used in `request.form[]` inside app.py. The form uses `method='POST'` and `action='/predict'`. All inputs are `type='number'` so only numeric values can be entered.

### The 9 Input Fields (matching model features):

- `age_first_funding_year` — Years after founding when first funding was received
- `age_last_funding_year` — Years after founding when last funding was received
- `age_first_milestone_year` — Years after founding when first milestone was achieved
- `age_last_milestone_year` — Years after founding when last milestone was achieved
- `relationships` — Total number of business relationships
- `funding_rounds` — Number of funding rounds completed
- `funding_total_usd` — Total funding in US Dollars
- `milestones` — Total number of milestones achieved
- `avg_participants` — Average investors per funding round

### **predict.html — Form Structure (Key Code):**

```
<!DOCTYPE html>  
  
<html lang="en">  
  
<head>  
    <meta charset="UTF-8">  
    <title>Predict - Prosperity Prognosticator</title>  
    <style>  
        body { font-family: 'Segoe UI', sans-serif;  
                background: linear-gradient(135deg, #0d1b2a, #1b2a4a);  
        }</style>
```

```
        min-height:100vh; color:#fff; margin:0; }
```

```
.form-card {
```

```
    max-width:820px; margin:40px auto;
```

```
    background: rgba(255,255,255,0.05);
```

```
    border: 1px solid rgba(255,255,255,0.12);
```

```
    border-radius:24px; padding:48px;
```

```
}
```

```
.form-grid { display:grid; grid-template-columns:1fr 1fr;
```

```
gap:22px; }
```

```
label { font-size:13px; color:#9aafc4; display:block; margin-
```

```
bottom:8px; }
```

```
input {
```

```
    width:100%; background:rgba(255,255,255,0.07);
```

```
    border:1px solid rgba(255,255,255,0.15);
```

```
    border-radius:12px; padding:14px 18px;
```

```
    color:#fff; font-size:15px;
```

```
}
```

```
.submit-btn {
```

```
    background: linear-gradient(135deg,#ffa600,#ff7b00);
```

```
    color:#0d1b2a; padding:16px 60px;
```

```
    border-radius:30px; border:none;
```

```
    font-size:16px; font-weight:700; cursor:pointer;
```

```
}
```

```
</style>
```

```
</head>
```

```
<body>
```

```
    <div class="form-card">
```

```
        <h1>Startup Success Predictor</h1>
```

```
        <form action="/predict" method="POST">
```

```
            <div class="form-grid">
```

```
                <div>
```

```
                    <label>Age at First Funding (years)</label>
```

```
                    <input type="number" name="age_first_funding_year"
```

```
                        step="0.1" placeholder="e.g. 1.5" required>
```

```
                </div>
```

```
                <div>
```

```
                    <label>Age at Last Funding (years)</label>
```

```
                    <input type="number" name="age_last_funding_year"
```

```
                        step="0.1" placeholder="e.g. 4.2" required>
```

```
                </div>
```

```
                <div>
```

```

        <label>Age at First Milestone (years)</label>
        <input type="number" name="age_first_milestone_year"
               step="0.1" placeholder="e.g. 0.5" required>
    </div>
    <div>
        <label>Age at Last Milestone (years)</label>
        <input type="number" name="age_last_milestone_year"
               step="0.1" placeholder="e.g. 3.0" required>
    </div>
    <div>
        <label>Number of Relationships</label>
        <input type="number" name="relationships"
               step="1" placeholder="e.g. 5" required>
    </div>
    <div>
        <label>Funding Rounds</label>
        <input type="number" name="funding_rounds"
               step="1" placeholder="e.g. 3" required>
    </div>
    <div>
        <label>Total Funding (USD)</label>
        <input type="number" name="funding_total_usd"
               step="1000" placeholder="e.g. 500000" required>
    </div>
    <div>
        <label>Total Milestones</label>
        <input type="number" name="milestones"
               step="1" placeholder="e.g. 4" required>
    </div>
    <div>
        <label>Avg. Participants per Round</label>
        <input type="number" name="avg_participants"
               step="0.1" placeholder="e.g. 2.5" required>
    </div>
</div>
<div style="text-align:center; margin-top:30px;">
    <button type="submit" class="submit-btn">
        Predict My Startup's Future
    </button>
</div>
</form>

```

```
</div>
</body>
</html>
```

## Activity 3: Building submit.html (Result Page)

The submit/result page displays the model's prediction. It uses Flask's Jinja2 templating with an if/else block to show different content depending on whether the prediction is 'Acquired' (green success state) or 'Closed' (red warning state). The result is passed as the variable prediction\_text from app.py.

Jinja2 templating: Flask uses Jinja2 to embed Python logic inside HTML.  
{{ prediction\_text }} displays the variable value. {% if prediction\_text == 'Acquired' %} checks the condition. {% else %} handles the other case. {% endif %} closes the block.

### submit.html — Key Code (Jinja2 template logic):

```
<!DOCTYPE html>
<html lang="en">
<head>
    <meta charset="UTF-8">
    <title>Prediction Result</title>
    <style>
        body { font-family:'Segoe UI', sans-serif;
               background:linear-gradient(135deg, #0d1b2a, #1b2a4a);
               min-height:100vh; color:#fff;
               display:flex; align-items:center;
               justify-content:center; margin:0; }

        .result-card {
            background:rgba(255,255,255,0.05);
            border:1px solid rgba(255,255,255,0.12);
            border-radius:28px; padding:60px;
            max-width:680px; text-align:center;
        }

        .icon-acquired { font-size:80px; color:#00c851; }
        .icon-closed   { font-size:80px; color:#ff4444; }
        .btn { padding:14px 38px; border-radius:28px;
               text-decoration:none; font-weight:700; }

        .btn-green { background:linear-gradient(135deg, #00c851, #007e33);
                     color:#fff; }

        .btn-red   { background:linear-gradient(135deg, #ff4444, #cc0000); }
```

```

        color:#fff; }

.btn-back { background:rgba(255,255,255,0.08); color:#fff;
            border:1px solid rgba(255,255,255,0.2); }

</style>
</head>
<body>

<div class="result-card">

    {% if prediction_text == 'Acquired' %}

        <div class="icon-acquired"> </div>
        <h2 style="color:#00c851; margin:20px 0;">
            Your Startup is Likely to be Acquired!
        </h2>
        <p>Strong indicators of success detected.</p>
        <p><strong>Prediction: ACQUIRED (Status = 1)</strong></p>
        <br>
        <a href="/predict" class="btn btn-green">Predict Another</a>
        <a href="/" class="btn btn-back">Home</a>

    {% else %}

        <div class="icon-closed"> </div>
        <h2 style="color:#ff4444; margin:20px 0;">
            Your Startup is at Risk of Closing
        </h2>
        <p>Consider improving funding and milestone metrics.</p>
        <p><strong>Prediction: CLOSED (Status = 0)</strong></p>
        <br>
        <a href="/predict" class="btn btn-red">Try Again</a>
        <a href="/" class="btn btn-back">Home</a>

    {% endif %}

</div>
</body>
</html>
```

#### **Activity 4: Building app.py (Flask Server)**

app.py is the heart of the web application. It starts the Flask server, defines URL routes, loads the trained model, receives form data from predict.html, runs the model prediction, and sends the result to submit.html.

Flask routing: `@app.route('/')` connects a URL path to a Python function. When a user visits `http://127.0.0.1:5000/`, Flask calls the `home()` function. When they submit the form (POST to `/predict`), Flask calls the `predict()` function.

### Step 1 — Import Libraries and Load Model:

#### **Code:**

```
from flask import Flask, request, render_template
import joblib
import numpy as np

# Initialize Flask application
app = Flask(__name__)

# Load the saved Random Forest model
# This file must exist in the same folder as app.py
model = joblib.load('random_forest_model.pkl')
```

### Step 2 — Home Page Route:

#### **Code:**

```
@app.route('/')
def home():
    """Render the home/landing page."""
    return render_template('home.html')
```

### Step 3 — Predict Route (GET + POST):

#### **Code:**

```
@app.route('/predict', methods=['GET', 'POST'])
def predict():
    if request.method == 'POST':
        # Retrieve all 9 form values in the correct order
        features = [
            float(request.form['age_first_funding_year']),
            float(request.form['age_last_funding_year']),
            float(request.form['age_first_milestone_year']),
            float(request.form['age_last_milestone_year']),
            float(request.form['relationships']),
            float(request.form['funding_rounds']),
            float(request.form['funding_total_usd']),
```

```

        float(request.form['milestones']),
        float(request.form['avg_participants']),
    ]
# Reshape for model input (1 row, 9 features)
final_input = np.array(features).reshape(1, -1)
# Predict: 1 = Acquired, 0 = Closed
prediction = model.predict(final_input)
result = 'Acquired' if prediction[0] == 1 else 'Closed'
return render_template('submit.html', prediction_text=result)

# GET request: show the empty form
return render_template('predict.html')

```

#### Step 4 — Main Function:

##### **Code:**

```

if __name__ == '__main__':
    app.run(debug=True)

```

#### **Activity 5: How to Run the Application**

Once all files are created and the model is saved, follow these steps exactly to launch the web application and access it in your browser.

Before running: Make sure random\_forest\_model.pkl is in the SAME folder as app.py. If it is not there, re-run Milestone 7 (the notebook) to generate it first.

#### Step 1 — Open Anaconda Prompt as Administrator:

- Click the Windows Start Menu
- Search for 'Anaconda Prompt'
- Right-click it and select 'Run as Administrator'

#### Step 2 — Navigate to Project Folder:

##### **Command:**

```

# Replace the path below with your actual project folder path
cd C:\Users\YourName\ProsperityPrognosticator

# Verify you are in the right folder
dir

```

#### Step 3 — Install Flask (if not already installed):

##### **Command:**

```
pip install flask  
pip install joblib
```

#### Step 4 — Run the Flask Application:

##### **Command:**

```
python app.py
```

After running this command, you should see output similar to:

```
* Serving Flask app 'app'  
* Debug mode: on  
* Running on http://127.0.0.1:5000  
* Press CTRL+C to quit
```

#### Step 5 — Open the Application in Browser:

- Open Google Chrome, Firefox, or any browser
- In the address bar, type exactly: `http://127.0.0.1:5000`
- Press Enter — the home page should load
- Click 'Predict Now' to go to the prediction form
- Fill in all 9 fields and click 'Predict My Startup's Future'
- The result page shows 'Acquired' or 'Closed' with a visual indicator

#### Step 6 — Stop the Application:

- Go back to the Anaconda Prompt window
- Press `Ctrl + C` on the keyboard to stop the Flask server
- The prompt will return — the app is now stopped

# AI Startup Success Predictor

Enter your startup metrics to get an AI-powered success prediction

Age at First Funding (Years)

1

Age at Last Funding (Years)

2

Age at First Milestone (Years)

2

Age at Last Milestone (Years)

1

Number of Relationships

3

Funding Rounds

2

Total Funding (USD)

e.g., 5000000

I



Number of Milestones

e.g., 8

Average Participants

e.g., 12

ACQUIRED ✓



## The Startup Is: Acquired

AI-Powered Startup Success Classification

Model Confidence: 92.7% • Powered by Machine Learning

**4.2**

SUCCESS SCORE

**87%**

PREDICTION ACCURACY

**1,247**

SIMILAR STARTUPS

CLOSED ⚠



# The Startup Is: Closed

AI-Powered Startup Success Classification

Model Confidence: 92.7% • Powered by Machine Learning

**4.2**

SUCCESS SCORE

**87%**

PREDICTION ACCURACY

**1,247**

SIMILAR STARTUPS

## Project Summary

The table below summarises all milestones, key activities, and outputs produced across the entire Prosperity Prognosticator project:

Milestone	Key Activities	Output
<b>1 – Data Collection</b>	Download startup1.csv, import libraries, load dataset, inspect head/tail/info	Raw DataFrame loaded
<b>2 – Data Information</b>	State grouping pie chart, separate numeric/categorical, convert status, drop labels	Cleaned target variable

<b>3 – Visualisation</b>	Statistical describe(), value counts, missing values, correlation heatmap, scatter & box plots	EDA complete
<b>4 – Univariate Analysis</b>	Class counts, pie chart (acquired vs closed), countplot per category	Class distribution understood
<b>5 – Bivariate Analysis</b>	Success rate per category, total funding per category, startup count per state	Feature-outcome insights
<b>6 – Pre-processing</b>	Remove duplicates, drop negative values, outlier analysis, drop irrelevant columns, train-test split	Clean X_train / X_test
<b>7 – Model Building</b>	Random Forest training, confusion matrix, classification report, ROC-AUC, Precision-Recall, save model	random_forest_model.pkl